

Linear Regression - Association of Tennis Professionals data

May 10, 2024

```
[8]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# load and investigate the data here:
df = pd.read_csv('./tennis_stats.csv')

# perform exploratory analysis here:
df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1721 entries, 0 to 1720

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	Player	1721 non-null	object
1	Year	1721 non-null	int64
2	FirstServe	1721 non-null	float64
3	FirstServePointsWon	1721 non-null	float64
4	FirstServeReturnPointsWon	1721 non-null	float64
5	SecondServePointsWon	1721 non-null	float64
6	SecondServeReturnPointsWon	1721 non-null	float64
7	Aces	1721 non-null	int64
8	BreakPointsConverted	1721 non-null	float64
9	BreakPointsFaced	1721 non-null	int64
10	BreakPointsOpportunities	1721 non-null	int64
11	BreakPointsSaved	1721 non-null	float64
12	DoubleFaults	1721 non-null	int64
13	ReturnGamesPlayed	1721 non-null	int64
14	ReturnGamesWon	1721 non-null	float64
15	ReturnPointsWon	1721 non-null	float64
16	ServiceGamesPlayed	1721 non-null	int64
17	ServiceGamesWon	1721 non-null	float64
18	TotalPointsWon	1721 non-null	float64
19	TotalServicePointsWon	1721 non-null	float64
20	Wins	1721 non-null	int64
21	Losses	1721 non-null	int64

```

22 Winnings          1721 non-null    int64
23 Ranking           1721 non-null    int64
dtypes: float64(12), int64(11), object(1)
memory usage: 322.8+ KB

```

```
[11]: df.describe()
```

```

[11]:
      count      Year  FirstServe  FirstServePointsWon  \
count  1721.000000  1721.000000          1721.000000
mean    2013.646717    0.598053            0.680738
std         2.488018    0.054533            0.070422
min     2009.000000    0.360000            0.270000
25%     2012.000000    0.570000            0.650000
50%     2014.000000    0.600000            0.690000
75%     2016.000000    0.630000            0.720000
max     2017.000000    0.880000            0.890000

      count  FirstServeReturnPointsWon  SecondServePointsWon  \
count          1721.000000          1721.000000
mean              0.261673            0.479733
std              0.056639            0.066902
min              0.000000            0.060000
25%              0.240000            0.460000
50%              0.270000            0.490000
75%              0.290000            0.520000
max              0.480000            0.920000

      count  SecondServeReturnPointsWon      Aces  BreakPointsConverted  \
count          1721.000000  1721.000000          1721.000000
mean              0.466432    97.105171            0.369407
std              0.068447   137.966077            0.162987
min              0.000000    0.000000            0.000000
25%              0.440000    7.000000            0.320000
50%              0.480000   34.000000            0.380000
75%              0.500000  140.000000            0.430000
max              0.750000  1185.000000            1.000000

      count  BreakPointsFaced  BreakPointsOpportunities  ...  ReturnGamesWon  \
count          1721.000000          1721.000000  ...          1721.000000
mean           112.003486          102.918071  ...            0.173823
std           119.247651          122.761670  ...            0.080880
min              1.000000           0.000000  ...            0.000000
25%           15.000000           9.000000  ...            0.130000
50%           55.000000          41.000000  ...            0.180000
75%          201.000000         172.000000  ...            0.220000
max          507.000000         573.000000  ...            0.560000

```

	ReturnPointsWon	ServiceGamesPlayed	ServiceGamesWon	TotalPointsWon \
count	1721.000000	1721.000000	1721.000000	1721.000000
mean	0.342208	197.650203	0.715590	0.473155
std	0.049369	221.208703	0.123287	0.037139
min	0.000000	0.000000	0.000000	0.220000
25%	0.320000	22.000000	0.670000	0.460000
50%	0.350000	86.000000	0.750000	0.480000
75%	0.370000	348.000000	0.790000	0.500000
max	0.510000	916.000000	1.000000	0.820000

	TotalServicePointsWon	Wins	Losses	Winnings \
count	1721.000000	1721.000000	1721.000000	1.721000e+03
mean	0.599245	7.876816	9.278908	2.344928e+05
std	0.057718	10.183716	8.996450	2.530537e+05
min	0.250000	0.000000	0.000000	1.080000e+02
25%	0.570000	0.000000	2.000000	4.931100e+04
50%	0.610000	3.000000	5.000000	1.252120e+05
75%	0.630000	13.000000	17.000000	3.500750e+05
max	0.820000	48.000000	36.000000	1.074562e+06

	Ranking
count	1721.000000
mean	269.610691
std	277.341947
min	3.000000
25%	83.000000
50%	166.000000
75%	333.000000
max	1443.000000

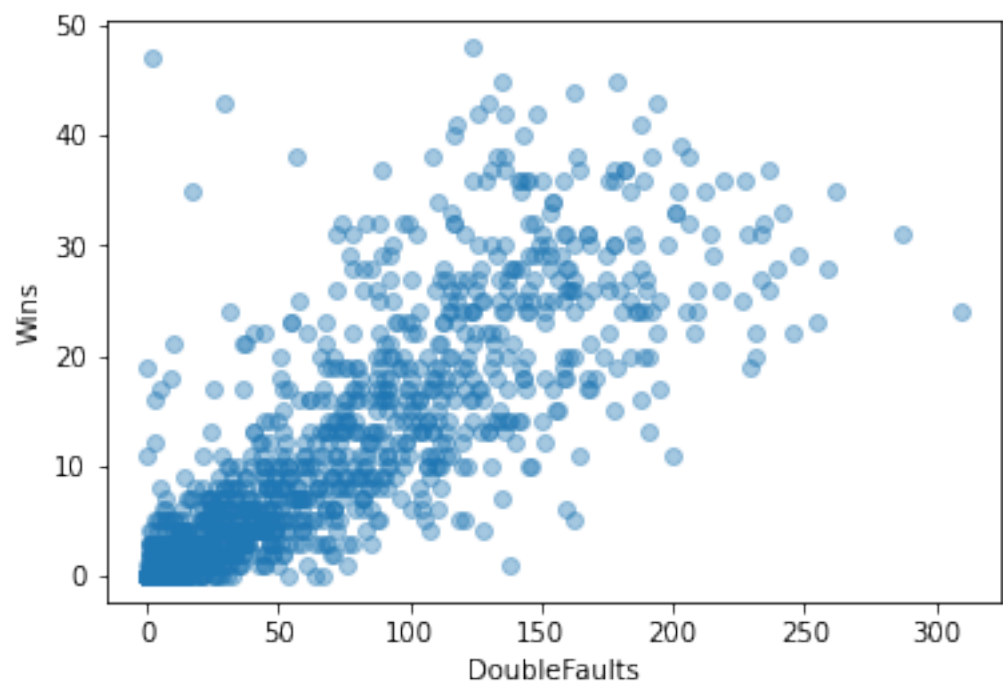

[8 rows x 23 columns]



```
[34]: # Explore features visually
plt.plot(df['Aces'], df['Wins'], 'o', alpha=0.4) # good feature
plt.xlabel("Aces")
plt.ylabel("Wins")
plt.show()
plt.plot(df['DoubleFaults'], df['Wins'], 'o', alpha=0.4) # good feature
plt.xlabel("DoubleFaults")
plt.ylabel("Wins")
plt.show()
plt.plot(df['ServiceGamesWon'], df['Wins'], 'o', alpha=0.4)
plt.xlabel("ServiceGamesWon")
plt.ylabel("Wins")
plt.show()
plt.plot(df['ReturnPointsWon'], df['Wins'], 'o', alpha=0.4)
plt.xlabel("ReturnPointsWon")
```

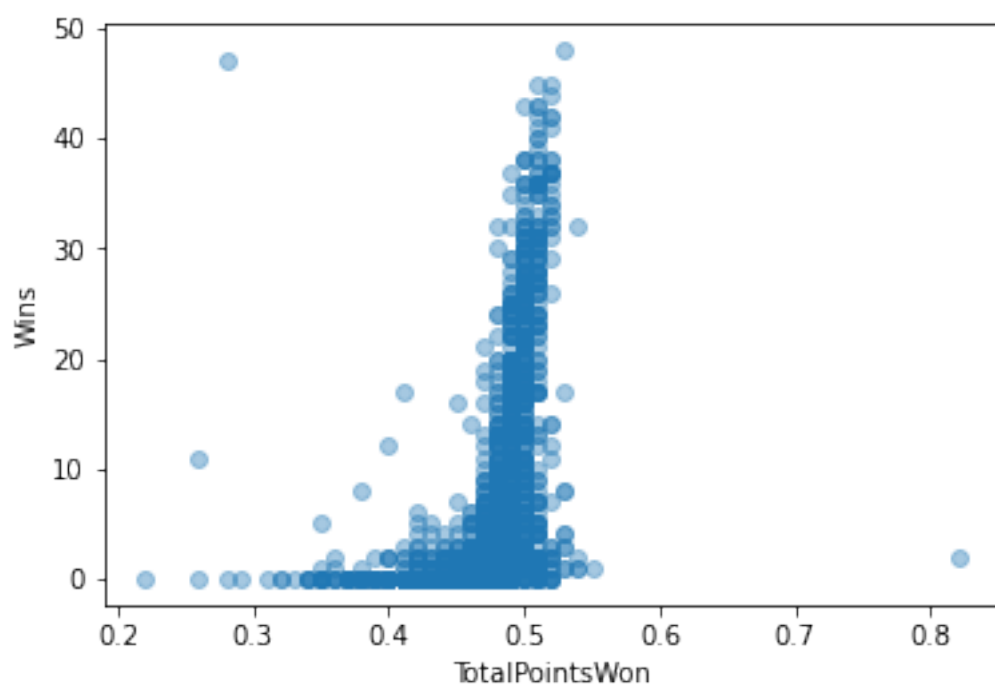

```

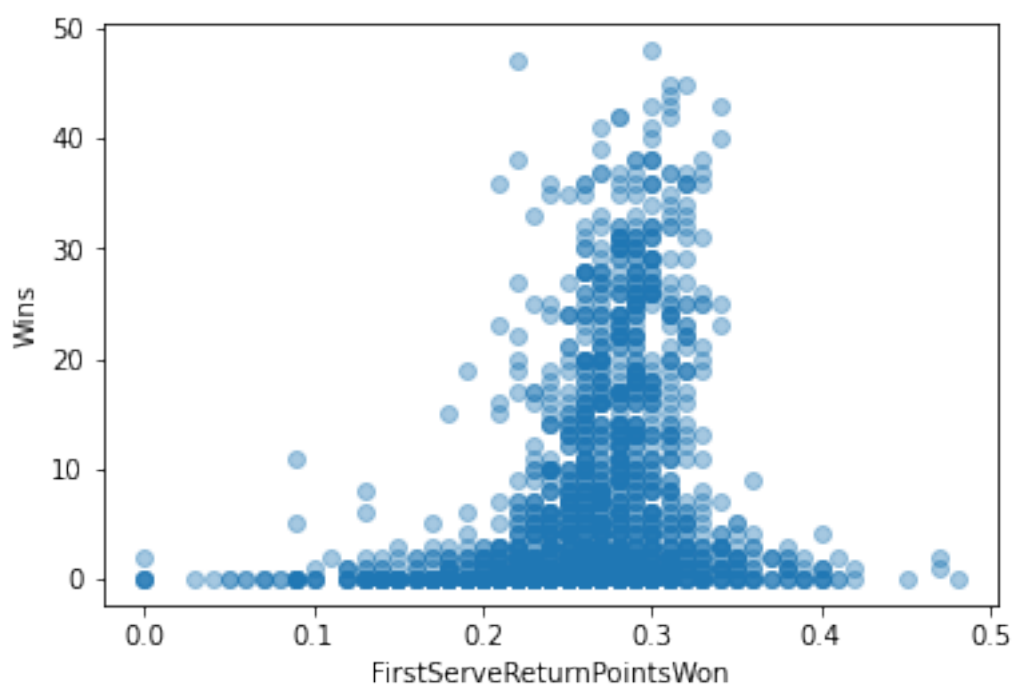

plt.ylabel("Wins")
plt.show()
plt.plot(df['ReturnGamesWon'], df['Wins'], 'o', alpha=0.4)
plt.xlabel("ReturnGamesWon")
plt.ylabel("Wins")
plt.show()
plt.plot(df['TotalPointsWon'], df['Wins'], 'o', alpha=0.4)
plt.xlabel("TotalPointsWon")
plt.ylabel("Wins")
plt.show()
plt.plot(df['FirstServe'], df['Wins'], 'o', alpha=0.4)
plt.xlabel("FirstServe")
plt.ylabel("Wins")
plt.show()
plt.plot(df['FirstServeReturnPointsWon'], df['Wins'], 'o', alpha=0.4)
plt.xlabel("FirstServeReturnPointsWon")
plt.ylabel("Wins")
plt.show()
plt.plot(df['BreakPointsFaced'], df['Wins'], 'o', alpha=0.4) #good feature
plt.xlabel("BreakPointsFaced")
plt.ylabel("Wins")
plt.show()
plt.plot(df['BreakPointsConverted'], df['Wins'], 'o', alpha=0.4) #good feature
plt.xlabel("BreakPointsConverted")
plt.ylabel("Wins")
plt.show()
plt.plot(df['BreakPointsOpportunities'], df['Wins'], 'o', alpha=0.4) #good
    ↪ feature
plt.xlabel("BreakPointsOpportunities")
plt.ylabel("Wins")
plt.show()

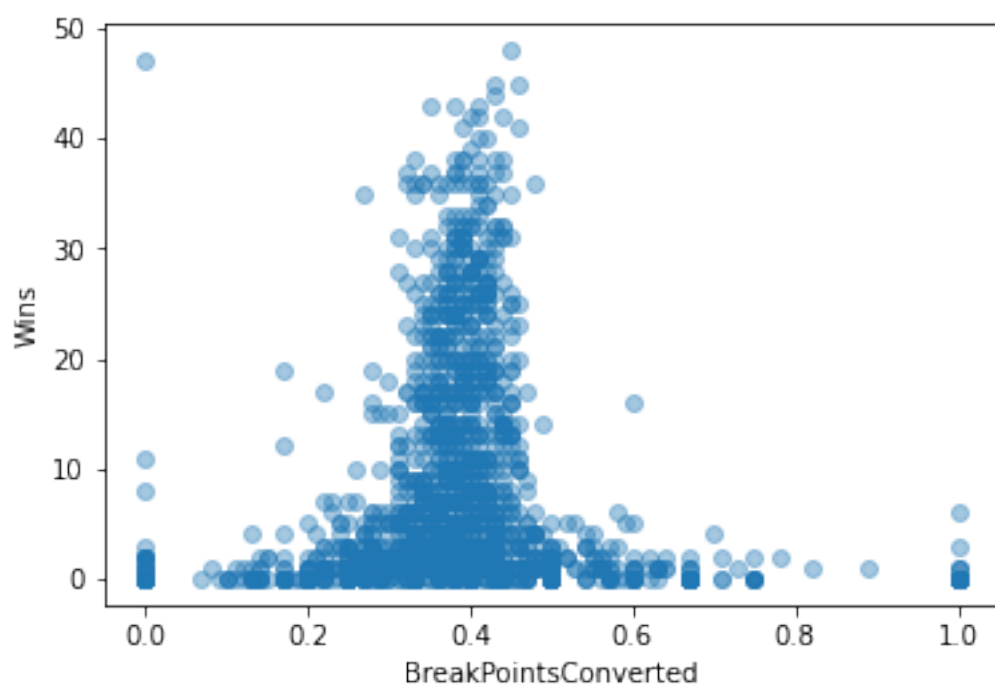

```














```
[35]: # create train/test data split
features = df[['BreakPointsOpportunities']]
labels = df['Wins']


x_train, x_test, y_train, y_test = train_test_split(features, labels,
↳train_size=0.8)

[41]: # Single feature linear regression
SingleLR = LinearRegression()
SingleLR.fit(x_train, y_train)
print("The line for BreakPoints Opportunities vs Wins has a slope of %.4f and
↳an intercept of %.4f" % (SingleLR.coef_, SingleLR.intercept_))
R_squared = SingleLR.score(x_test, y_test)
print("This model has an R^2 value of %.2f" % R_squared)
```

The line for BreakPoints Opportunities vs Wins has a slope of 0.0760 and an intercept of -0.0094
This model has an R^2 value of 0.86


```
[45]: prediction = SingleLR.predict(x_test)
plt.scatter(y_test, prediction, alpha=0.4)
plt.xlabel("Actual Wins")
plt.ylabel("Predicted Wins")
plt.title("Predicted vs Actual wins based on BreakPointsOpportunities")
```

```
plt.show()
```



```
[51]: # New single feature linear regression
features = df[['Aces']]
x_train, x_test, y_train, y_test = train_test_split(features, labels,
    ↪ train_size=0.8)
SingleLR = LinearRegression()
SingleLR.fit(x_train, y_train)
print("The line for Aces vs Wins has a slope of %.4f and an intercept of %.4f"
    ↪ % (SingleLR.coef_, SingleLR.intercept_))
R_squared = SingleLR.score(x_test, y_test)
print("This model has an R^2 value of %.2f" % R_squared)
prediction = SingleLR.predict(x_test)
plt.scatter(y_test, prediction, alpha=0.4)
plt.xlabel("Actual Wins")
plt.ylabel("Predicted Wins")
plt.title("Predicted vs Actual wins based on Aces")
plt.show()
```

The line for Aces vs Wins has a slope of 0.0624 and an intercept of 1.8811
This model has an R² value of 0.70



```
[54]: # Double feature linear regression
features = df[['BreakPointsOpportunities', 'BreakPointsFaced']]
labels = df[['Winnings']]
DoubleLR = LinearRegression()
x_train, x_test, y_train, y_test = train_test_split(features, labels,
    ↪train_size=0.8)
DoubleLR.fit(x_train, y_train)
r2 = DoubleLR.score(x_test, y_test)
print("R^2 for BreakPointsOpportunities and BreakPointsFaced vs Winnings is %.
    ↪2f" % r2)
features = df[['Aces', 'DoubleFaults']]
DoubleLR = LinearRegression()
x_train, x_test, y_train, y_test = train_test_split(features, labels,
    ↪train_size=0.8)
DoubleLR.fit(x_train, y_train)
r2 = DoubleLR.score(x_test, y_test)
print("R^2 for Aces and DoubleFaults vs Winnings is %.2f" % r2)
```

R² for BreakPointsOpportunities and BreakPointsFaced vs Winnings is 0.80

R² for Aces and DoubleFaults vs Winnings is 0.72

```
[56]: print("BreakPoints Opportunities and BreakPoints Faced are better predictors,
    ↪for Winnings than Aces and DoubleFaults")
```

BreakPoints Opportunities and BreakPoints Faced are better predictors for Winnings than Aces and DoubleFaults

```
[61]: # multiple linear regression
from sklearn.preprocessing import StandardScaler
mlr = LinearRegression()
features = df.drop(columns=['Player', 'Wins', 'Losses', 'Winnings', 'Ranking'])
labels = df['Winnings']
features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1721 entries, 0 to 1720
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Year                                1721 non-null   int64
1   FirstServe                          1721 non-null   float64
2   FirstServePointsWon                 1721 non-null   float64
3   FirstServeReturnPointsWon           1721 non-null   float64
4   SecondServePointsWon                1721 non-null   float64
5   SecondServeReturnPointsWon          1721 non-null   float64
6   Aces                                1721 non-null   int64
7   BreakPointsConverted                1721 non-null   float64
8   BreakPointsFaced                    1721 non-null   int64
9   BreakPointsOpportunities            1721 non-null   int64
10  BreakPointsSaved                     1721 non-null   float64
11  DoubleFaults                        1721 non-null   int64
12  ReturnGamesPlayed                   1721 non-null   int64
13  ReturnGamesWon                      1721 non-null   float64
14  ReturnPointsWon                     1721 non-null   float64
15  ServiceGamesPlayed                  1721 non-null   int64
16  ServiceGamesWon                     1721 non-null   float64
17  TotalPointsWon                      1721 non-null   float64
18  TotalServicePointsWon                1721 non-null   float64
dtypes: float64(12), int64(7)
memory usage: 255.6 KB
```

```
[65]: scaler = StandardScaler()
scaledFeatures = scaler.fit_transform(features)
x_train, x_test, y_train, y_test = train_test_split(scaledFeatures, labels,
    ↪train_size=0.8)
mlr.fit(x_train, y_train)
r2 = mlr.score(x_test, y_test)
print("Using a standardized form of all numerical features, the R^2 value for_
    ↪multiple linear regression predicting Winnings is %.2f" % r2)
```

Using a standardized form of all numerical features, the R² value for multiple linear regression predicting Winnings is 0.82

[]: